Here are the direct links to the best datasets and the official X algorithm source code:

---

1. Interaction & Engagement Data

This is the "gold mine" for training models that predict what a user will like or retweet.

- **RecSys Challenge 2020 (Twitter Dataset)**
  - **What it is:** A massive dataset of 200 million public engagements (Likes, Replies, Retweets) released by Twitter for researchers.
  - **Direct Access:** Often hosted on **Academic Torrents** due to its size (70GB+).

2. Social Graph & Network Data

Use these to build "Who to Follow" recommendations based on who is connected to whom.

- **Stanford SNAP Twitter Network**
  - **What it is:** A collection of "social circles" (lists) and ego networks. It's perfect for Graph Neural Networks (GNNs).
- **Twitter Friends (Kaggle)**
  - **What it is:** A smaller, cleaner dataset showing user IDs and the list of IDs they follow. Great for building a quick prototype.

3. Content & Sentiment Data

Use these to understand *what* is being said so you can recommend similar topics.

- **Sentiment140 (1.6 Million Tweets)**
  - **What it is:** The standard dataset for NLP on Twitter text.
- **Customer Support on Twitter**
  - **What it is:** Natural conversational data (replies and threads) which is helpful for understanding "engagement loops."

---

4. The Official Source Code

If you want to see the actual logic X uses to rank posts, they open-sourced it on GitHub.

- **X (Twitter) Recommendation Algorithm GitHub**
  - This repository contains the **Heavy Ranker**, **SimClusters** (community detection), and the **Graph Feature Service**.
- **The Machine Learning Models**
  - This specifically contains the Python/Torch code for the "For You" timeline ranking models.

Note: Because X's terms of service change frequently, some direct download buttons on Kaggle may require you to "Accept Terms" before the link works.

Would you like a sample Python code snippet to show you how to load the SNAP graph data and find "friends of friends"?